

Data Structures & Algorithms

**Week 8 - Priority Queues (Binary Heaps, Skew Heaps,
Applications)**

Subodh Sharma, Rahul Garg
{svs,rahulgarg}@iitd.ac.in

Priority Queues

- An **ADT** similar to a Queue or Stack but with a **caveat**
 - Each element has an associated **priority**
- **Motivation:**
 - Many application tasks running on OS, and you press ESC (or Ctrl-C)! What would you expect?
 - What would have happened if every task had the same priority?
- **Applications:**
 - Scheduling, Algorithmic efficiency (Spanning Trees, Shortest Paths etc.), Simulation Systems (Discrete Event Simulation, etc.), Network Traffic Mgmt. (routing pkts with different service reqs.), E-commerce, Load balancing, etc.

On Priorities

IF YOU DON'T TURN IN
AT LEAST ONE HOMEWORK
ASSIGNMENT, YOU'LL
FAIL THIS CLASS.

YEAH. BUT IF I CAN FAIL
THIS CLASS, THE GRADES
ON MY REPORT CARD WILL
BE IN ALPHABETICAL ORDER!



Courtesy: XKCD Comics

On Priorities

- Priorities help rank the elements in a Priority Queue with a **total order relation**
- **Total order relation:**
 - **Reflexive:** $a \leq a$
 - **Antisymmetric:** if $a_1 \leq a_2$ and $a_2 \leq a_1$, then $a_1 = a_2$
 - **Transitive:** if $a_1 \leq a_2$ and $a_2 \leq a_3$, then $a_1 \leq a_3$

Priority Queue: Model

- Supported operations: **Insert** and **DeleteMin**


- Various implementation of PQ:

- Using simple **linked lists**:

- Insertions at the front — $O(1)$
 - Deleting the minimum — $O(N)$


- Using **linked lists** that **remain sorted**:

- Insertions — $O(N)$
 - DeleteMin (from front) — $O(1)$




(Binary) Heap

- Heap is a **binary tree** that stores priority or priority-value pairs at its nodes
- Heaps have two important properties:
 - **Structure Property:** Heap is **completely filled** with the **exception of the last level**.
 - The last level is **left-filled**.
 - **Order Property:** Every node should be **smaller** than all of its descendants




Examples of Non Heaps

- Structure property violation



- Order property violation





Height of the Heap

- Suppose a heap of n nodes has height h
- Complete binary tree of height h has $2^{h+1} - 1$ nodes
- Hence $2^h - 1 < n \leq 2^{h+1} - 1$
- Thus, $h = \lfloor \log_2 n \rfloor$

Implementing Heaps



- **Observation:** Complete binary tree is **so regular** that it can be represented **by arrays** instead of pointers

```
int getParentIndex(int i) {  
    return (i - 1) / 2;  
}  
  
int getLeftChildIndex(int i) {  
    return 2 * i + 1;  
}  
  
int getRightChildIndex(int i) {  
    return 2 * i + 2;  
}
```



Implementing Heaps: Efficiency


- **Observation:** In binary representation, **multiplication by 2 is a left shift** and FMA instructions to multiply and add (adding 1 to the lowest bit)



Heap Insertion


- Insert 12
- The process of restoring order is called **Heapify**

```
void insert(int val) {  
    heap.push_back(val);  
    heapifyUp(heap.size() - 1);  
}  
  
void heapifyUp(int index) {  
    if (index == 0) return;  
  
    int parentIndex = getParentIndex(index);  
  
    if (heap[parentIndex] > heap[index]) {  
        swap(heap[parentIndex], heap[index]);  
        heapifyUp(parentIndex);  
    }  
}
```




Why is Insertion Correct?

- The only nodes whose contents change are the ones on the path
- Heap property may violate only for children of these nodes
- But the new contents of these nodes only have lower priority
- Thus, it is correct!




Another View of Heapify

- Heap order property violated at index 0
- The subtrees rooted at index 1 and 2 are valid heaps
 - This is an important point — Heapify would work only when this observation holds
- heapifyDown(0)
- **ToDo — Prove the correctness of HeapifyDown**



HeapifyDown

```
void heapifyDown(int index) {  
    int leftChild = getLeftChildIndex(index);  
    int rightChild = getRightChildIndex(index);  
  
    if (leftChild >= heap.size()) return; // No children  
  
    int minIndex = index;  
  
    if (heap[minIndex] > heap[leftChild]) {  
        minIndex = leftChild;  
    }  
  
    if (rightChild < heap.size() && heap[minIndex] > heap[rightChild]) {  
        minIndex = rightChild;  
    }  
  
    if (minIndex != index) {  
        swap(heap[minIndex], heap[index]);  
        heapifyDown(minIndex);  
    }  
}
```




DeleteMin Operation

- Remember that the minimum element is at the root of the heap
 - We can delete this and move one of its children to fill the space!
 - Empty location moves down the tree
 - Resulting tree **may not be left-filled**

DeleteMin Operation: Attempt 1


- Delete 13
- 16 moves up
- 19 moves up
- Not left-filled



DeleteMin Operation: Attempt 2

- Replace root element with the last element of the heap
- HeapifyDown(rootIndex)

```
void deleteMin() {  
    if (heap.empty()) {  
        std::cout << "Heap is empty!" << std::endl;  
        return;  
    }  
  
    heap[0] = heap.back();  
    heap.pop_back();  
  
    heapifyDown(0);  
}
```




Building Heap

- Simple method — Repeatedly call insert method
 - Time complexity: $\sum_{i=1}^n \log i = O(\log n!) = O(n \log n)$
- Better solution: We start from the bottom and move up
- All leaves are heaps (inductive construction)

```
void buildHeap(const std::vector<int> &arr) {  
    heap = arr;  
    int n = heap.size();  
  
    for (int i = n / 2 - 1; i >= 0; i--) {  
        heapifyDown(i);  
    }  
}
```

Building Heap

```
void buildHeap(const std::vector<int> &arr) {  
    heap = arr;  
    int n = heap.size();  
  
    for (int i = n / 2 - 1; i >= 0; i--) {  
        heapifyDown(i);  
    }  
}
```



Building Heap: Analysis

- Height of node : length of longest path from the node to leaf
- Height of tree: height of root
- Time for HeapifyX(i): $O(\text{height of the subtree rooted at } i)$
- Assume: $n = 2^k - 1$ (a complete binary tree — only help us simplify the analysis)

Building Heap: Analysis

- For the $n/2$ nodes of height 1: Heapify requires at most 1 swap
- For the $n/4$ nodes of height 2: Heapify requires at most 2 swaps
- For the $n/2^i$ nodes of height i : Heapify requires at most i swaps
- Total number of swaps: $\sum_{i=1}^{\log n} n \cdot i/2^i = O(n)$


Heap Sort

- Create a heap: $T(n) = O(n)$
- Do **DeleteMin** repeatedly till the heap becomes empty: $T(n) = O(n \log n)$
- Alternative strategy: No other space constraint, i.e., **in-place sort**
 - Do **DeleteMin** and move the deleted element to the end of the heap
 - Heapify the rest

In-Place Heap Sort

- NOTE: **Heap size is reduced by 1** after each such operation

```
void heapSort() {  
    int n = heap.size();  
    // Extract elements from heap one by one  
    for (int i = n - 1; i > 0; i--) {  
        // Move the root to the end  
        std::swap(heap[0], heap[i]);  
  
        // Heapify the reduced heap  
        heapifyDown(i, 0);  
    }  
}
```



- Time complexity: $O(n \log n)$

Runtime Analysis


- A heap of n nodes has height $O(\log n)$
- Insertion (heapifyUp along a path) — at most $O(\log n)$ steps
- HeapifyDown — $O(\log n)$
 - An element may be moved all the way to the last level
- DeleteMin — $O(\log n)$
- BuildHeap — $O(n)$
- HeapSort — $O(n \log n)$

Applications: The Selection Problem

- Problem: Find the k^{th} largest element in a list of n elements
- Algorithm 1A:
 - Read the elements in an array
 - Sort the array
 - Return the k^{th} indexed element from the sorted array
 - Time complexity: $O(n^2)$ with simple sorting; $O(n \log n)$ otherwise.


Applications: The Selection Problem

- Problem: Find the k^{th} largest element in a list of n elements
- Algorithm 1B:
 - Read **only** k elements in an array
 - Sort the array
 - The smallest is at k^{th} position. For the remaining elements:
 - Compare with the k^{th} element \rightarrow if the incoming element is larger then replace it with the k^{th} element
 - Time complexity: ?



Applications: The Selection Problem

- Changed Problem: Find the k^{th} **smallest** element in a list of n elements
- Algorithm2A:
 - Read elements in an array
 - Apply BuildHeap
 - Apply k DeleteMin operations
 - The last extracted element is our answer
 - Time complexity: $O(n + k \log n)$. Why?
 - What happens when $k = \lceil n/2 \rceil$ or when $k = O(n/\log n)$?




Applications: The Selection Problem

- Problem: Find the k^{th} **largest** element in a list of n elements

- Algorithm2B:

- Read **only** k elements in an array and build a minheap.
- New element is compared with the k^{th} largest
- If the new element is larger, it replaces the root
- At the end of the input, we return the root.
- Time complexity: $O(k + (n-k) \cdot \log k)$
- Can we do better? Quickselect $O(n)$ average time!



Standard algo, do quickselect on one side only.
 $T(n) = T(n/2) + n \rightarrow T(n) = 2n \rightarrow \text{done}$

Extra Reading

For those who want to challenge themselves

- Skew Heaps (efficient merge operations) # (Amortize Time Complexity Analysis)
- Binomial Queues ~~XX~~
- Fibonacci Heaps
- Find the median of an array efficiently # Quick select type only
- Understand Quick-select ~~XX~~